



Prof. Dr. Christoph Pflaum
Florian Schornbaum
Dominik Bartuschat

Winter Term
2014/2015

Simulation and Scientific Computing Assignment 1

Exercise 1 (*Matrix-Matrix Multiplication*)

Organization Details

1. **Your Tutors:** Your tutors for the exercises are Florian Schornbaum and Dominik Bartuschat. You can find their offices and email addresses on the official web page of the chair for system simulation:

<https://www10.informatik.uni-erlangen.de>

2. **Programming exercises:** The programming language for the implementation of the programming exercises is C++.
3. **Assignment sheets:** The assignment sheets and further information concerning this lecture are available on the web page for this lecture:

<https://www10.informatik.uni-erlangen.de/Teaching/Courses/WS2014/SiWiR/>.

4. **Team work:** In order to encourage team work, we require you to form teams of three people (if possible) submitting their solution together. The team formation is carried out in individual responsibility of the students. In front of the computer exercise room 00.131 at the chair for System Simulation a corresponding list is available for sign up. Each team member must sign up *by 31.10.2014 at the latest*.
5. **Credits:** Each assignment will be awarded with either a green, yellow, or red point. Earning one green and one red point is equivalent to earning two yellow points. In order to successfully participate in the exercise classes, you have to acquire at least three yellow and one green point. This also means that you are allowed to receive a red point only once!
6. **Tutorial:** In the context of the SiWiR tutorial we have planned several presentation sessions. Each participant is required to give a presentation and to attend every presentation session in order to successfully participate in the tutorial.
7. **Plagiarism:** It should go without saying that every team must provide us with their own solution, programmed by themselves. If a team ignores this, they will receive a red point for their assignment.

8. **Remote Access:** Remote access to our computers is only possible via the *secure shell (ssh)*. For the login, you first have to either access a computer from the university network, or log in via VPN, from where you can access the LSS CIP pool machines $i10cip\{1 \dots 12\}$. In our computer exercise room, the same rules as in the CIP pool of the computer science department apply:

<https://www.cip.informatik.uni-erlangen.de/cipPools/rules.en.html>.

Tasks

1. Your task is to implement a matrix-matrix multiplication $C = A \cdot B$, where A is a $M \times K$ matrix, B is a $K \times N$ matrix, and C is a $M \times N$ matrix.

Feel free to use every known algorithm and programming technique to decrease the single-core runtime of your program as long as you adhere to the following guidelines:

- All three matrices are represented as a linearized one-dimensional array with adjacent elements.
 - All three matrices are passed to your multiplication routine in a row-major format. Meaning, it is not allowed to store one of the input matrices in a transposed layout. However, the computation of the transpose is allowed to be a part of the multiplication routine itself such that the matrix is transposed after the start of the time measurement.
 - Make sure you use double precision floating-point operations for your multiplication.
 - The use of threads is prohibited.
2. You should use *likwid* (Like I knew what I am doing - <http://code.google.com/p/likwid/>) lightweight performance tools to measure the performance of your program and to gather information about the following events:
- L2 bandwidth,
 - L2 miss rate,
 - double precision FLOPS.

Important Note: Make sure that only your program is currently executing in order to measure reliable data!

A documentation for *likwid* can be found on the project's web page. A short introduction for the necessary *likwid* commands will be given in the exercise classes.

For your performance and *likwid* measurements, you have to use one of the computers in the LSS CIP pool. On the web page of this lecture, an example Makefile can be found that shows you how to compile your program that uses the *likwid* tool's marker API, which allows you to measure named regions in the code.

3. Your program should be able to be executed in the following form:

```
./matmult A.in B.in C.out,
```

where A.in and B.in are two files containing the two matrices which are multiplied and C.out is the output file for the resulting matrix. Use the following file format: the first line contains the number of rows and the number of columns of the matrix. From the second line on, each line

contains exactly one element in the row by row order $x_{11}, x_{12}, \dots, x_{1n}, x_{21}, \dots, x_{2n}, x_{31}, \dots, x_{mn}$. The beginning of a file for a 10×20 matrix might for example look like this:

```
10 20
0.1892
0.2783
0.4657
...
```

4. You must use at least two optimization techniques from the lecture or exercises.
5. Compare your calculation results to the reference files provided on the course website together with the corresponding input files! Make sure your implementation also works with non-square matrices.
6. Please hand in your solution to this exercise until Tuesday, November 4, 2014, 10:00 am. Make sure the following requirements are met:
 - (a) The program must be compilable with a Makefile. Your program must compile successfully when calling “make”.
 - (b) The program must be callable as specified above.
 - (c) The program must compile without errors or warnings on our computers in the LSS CIP pool with the following g++ compiler flags (you can make use of C++11 features):

```
-std=c++0x -Wall -Wextra -Wshadow -Werror -O3 -DNDEBUG
```

If you want to, you can add additional compiler flags.

The reference compiler is GCC version 4.9 on the CIP Pool computers. You can check the version by typing “g++ -version”. You can load GCC 4.9 by typing “module load gcc/4.9.0”.

- (d) The program must be compiled with *likwid* support by default and must output the elapsed wall clock time the actual computation of the multiplication takes. For *likwid* to work, you may have to load the corresponding module by typing “module load likwid”. In order to determine the time it takes for your algorithm to calculate the matrix-matrix multiplication, you must use the Timer class (“Timer.h”) provided on the course website. You only have to measure the time of your matrix-matrix multiplication, you do not have to measure the time for loading the matrices from disk and storing the result back to disk.
- (e) The solution should contain well commented source files, a fitting Makefile that satisfies all the conditions specified above, instructions how to use your program (e.g. in form of a README file), and a PDF file with performance graphs. When submitting the solution, remove all temporary files from your project directory and pack it using the following command:

```
tar -cjf ex01_groupXX.tar.bz2 ex01_groupXX/
```

where XX stands for your group number and ex01_groupXX/ contains your solution.

Then send the archive to the course mailing list:

```
siwir@i10.informatik.uni-erlangen.de
```

Performance challenge

We will test your code and measure the performance. Every code is compiled with the g++ compiler version 4.9.0. The program is run several times on one of the i10cip1...12 machines. We will use the problem size 2048^2 when measuring the performance.

The winning team will get promoted from a red to a yellow or from a yellow to a green point for one of their 4 assignments!

Beat Your Tutors!

The ultimate performance challenge: If you are the winning team of the performance challenge and on top of that can beat the time of “/software/opensuse/siwir/matmult”, you earn the right to present your program in the tutorial. This means you don’t have to become acquainted with another topic for the tutorial! Additionally, there will be another reward that will be disclosed once your program was demonstrated to beat our implementation.

Credits

In this assignment, points are awarded in the following way:

1. You are guaranteed to receive at least a yellow point if your program correctly performs the above tasks and fulfills all of the above requirements. The correctness will be assessed by checking whether your program matches the reference outputs. Make sure that your program also works with non-square matrices! Submissions with compile errors will lead to a red point! The computers in the LSS CIP pool act as reference environments.
2. A green point is awarded if additionally you provide measurements and corresponding graphs that show the effect of the used optimization techniques on runtime, double precision FLOPS, and L2 bandwidth and cache misses for the matrix-matrix multiplication. Perform the measurements for matrix sizes of 32^2 , 64^2 , 128^2 , 256^2 , 512^2 , 1024^2 , and 2048^2 . Plot well-labeled diagrams, one for L2 bandwidth, one for the L2 cache miss rate, one for the double precision FLOPS, and one for runtime, with the matrix size in \log_2 scale on the abscissa. Show in each diagram graphs for the standard implementation (without optimization), each single optimization, all optimizations combined, and the ATLAS implementation of the BLAS library function *dgemm*. Print all these diagrams in a pdf file, together with a description of information you can derive from the graphs (e.g. salient points) and with a comparison of the different implementations. Here you are encouraged to use information gained while optimizing the code with the help of the *likwid* tool.